

Neural system identification and control for Formula Student Driverless Cars

Final presentation

Tudor Oancea

EPFL Racing Team

June 2023

Goals:

- Produce a system model that is more precise than the first-principles models currently used by the EPFL RT Driverless team to create more lightweight and realistic (MiL) tests.
- Devise a control scheme that approximated well the existing MPC scheme while having a lighter computational and memory footprint.

Goals:

- Produce a system model that is more precise than the first-principles models currently used by the EPFL RT Driverless team to create more lightweight and realistic (MiL) tests.
- Devise a control scheme that approximated well the existing MPC scheme while having a lighter computational and memory footprint.


Introduction



Figure: Our test platform: the Formula Student Driverless Simulator (FSDS)

- 1 System identification
 - Bicycle model
 - Neural Ordinary Differential Equations
 - Dataset creation
 - Numerical experiments
- 2 Control
 - Original NMPC controller
 - Differentiable Predictive Control
 - Dataset creation
 - Numerical experiments
- 3 Conclusion & Outlook

Bicycle model - Notations



$$\tilde{x} = (X, Y, \varphi, v_x, v_y, r)$$
$$\tilde{u} = (T, \delta)$$

Bicycle model - Equations

Dynamic bicycle model (Dyn6)

$$\dot{X} = v_x \cos(\varphi) - v_y \sin(\varphi)$$

$$\dot{Y} = v_x \sin(\varphi) + v_y \cos(\varphi)$$

$$\dot{\varphi} = r$$

$$\dot{v}_x = \frac{1}{m}(F_x - F_{y,F} \sin(\delta) + mv_y r)$$

$$\dot{v}_y = \frac{1}{m}(F_{y,R} + F_{y,F} \cos(\delta) - mv_x r)$$

$$\dot{r} = \frac{1}{I_z}(F_{y,F} l_F \cos(\delta) - F_{y,R} l_R)$$

$$\text{with } F_x = (C_{m1} - C_{m2} v_x) T - C_{r0} \tanh(C_{r3} v_x) - C_{r2} v_x^2$$

$$F_{y,j} = D_j \sin(C_j \arctan(B_j \alpha_j)), j \in \{R, F\}$$

Bicycle model - Equations

Kinematic bicycle model (Kin4)

$$\dot{X} = v \cos(\varphi + \beta)$$

$$\dot{Y} = v \sin(\varphi + \beta)$$

$$\dot{\varphi} = \frac{v}{l_R} \sin(\beta)$$

$$\dot{v} = \frac{F_x}{m} \cos(\beta)$$

$$\text{with } \beta = \arctan \left(\frac{l_R}{l_R + l_F} \tan(\delta) \right)$$

Neural Ordinary Differential Equations - Architecture

Graybox NODE model (NeuralDyn6)

$$\dot{\tilde{\mathbf{x}}} = \begin{pmatrix} v_x \cos(\varphi) - v_y \sin(\varphi) \\ v_x \sin(\varphi) + v_y \cos(\varphi) \\ r \\ \nu_\theta(v_x, v_y, r, T, \delta) \end{pmatrix}$$

$$\text{FFNN } \nu_\theta : \begin{matrix} \mathbb{R}^5 & \rightarrow & \mathbb{R}^3 \\ (v_x, v_y, r, T, \delta) & \mapsto & (\dot{v}_x, \dot{v}_y, \dot{r}) \end{matrix}$$

RK4 discretization: $\tilde{\mathbf{x}}^+ = \tilde{f}_1(\tilde{\mathbf{x}}, \tilde{\mathbf{u}}) \implies \tilde{\mathbf{x}}_{1:N_f} = \tilde{f}_{N_f}(\tilde{\mathbf{x}}_0, \tilde{\mathbf{u}}_{0:N_f-1})$

$$a_{i:j} := (a_i, a_{i+1}, \dots, a_j)$$

Neural Ordinary Differential Equations - Architecture

Graybox NODE model (NeuralDyn6)

$$\dot{\tilde{x}} = \begin{pmatrix} v_x \cos(\varphi) - v_y \sin(\varphi) \\ v_x \sin(\varphi) + v_y \cos(\varphi) \\ r \\ \nu_\theta(v_x, v_y, r, T, \delta) \end{pmatrix}$$

FFNN $\nu_\theta : \mathbb{R}^5 \rightarrow \mathbb{R}^3$
 $(v_x, v_y, r, T, \delta) \mapsto (\dot{v}_x, \dot{v}_y, \dot{r})$

RK4 discretization: $\tilde{x}^+ = \tilde{f}_1(\tilde{x}, \tilde{u}) \implies \tilde{x}_{1:N_f} = \tilde{f}_{N_f}(\tilde{x}_0, \tilde{u}_{0:N_f-1})$

$$a_{i:j} := (a_i, a_{i+1}, \dots, a_j)$$

Neural Ordinary Differential Equations - Problem

Sysid dataset

$$\tilde{\mathcal{D}} = \{(\tilde{x}_0^i, \tilde{u}_{0:N_f-1}^i, \tilde{x}_{1:N_f}^i)\}_{i=1,\dots,I}$$

Sysid problem

$$\begin{aligned} \min_{\theta} \quad & \frac{1}{IN_f} \sum_{i=1}^I \sum_{k=1}^{N_f} \|\tilde{x}_k^i - \tilde{x}_k^{\text{pred},i}\|_P^2 \\ \text{s.t.} \quad & \tilde{x}_{1:N_f}^{\text{pred},i} = \tilde{f}_{N_f}(\tilde{x}_0^i, \tilde{u}_{0:N_f-1}^i), \quad i = 1, \dots, I \end{aligned}$$

- Time-series collected at 20Hz in FSDS with a video-game controller
 \implies transcription in the form $\tilde{x}_0^i, \tilde{u}_{0:N_f-1}^i, \tilde{x}_{1:N_f}^i$ with a sliding window
- Crowdsourced data collection during the EPFL open days \implies high volume of data from a wide range of operating scenarios.
- Pre-processing: yaw signal φ made continuous by appropriately translating each value by multiple of 2π .

Dataset creation

- Time-series collected at 20Hz in FSDS with a video-game controller \implies transcription in the form $\tilde{x}_0^i, \tilde{u}_{0:N_f-1}^i, \tilde{x}_{1:N_f}^i$ with a sliding window
- Crowdsourced data collection during the EPFL open days \implies high volume of data from a wide range of operating scenarios.
- Pre-processing: yaw signal φ made continuous by appropriately translating each value by multiple of 2π .

Dataset creation

- Time-series collected at 20Hz in FSDS with a video-game controller \implies transcription in the form $\tilde{x}_0^i, \tilde{u}_{0:N_f-1}^i, \tilde{x}_{1:N_f}^i$ with a sliding window
- Crowdsourced data collection during the EPFL open days \implies high volume of data from a wide range of operating scenarios.
- Pre-processing: yaw signal φ made continuous by appropriately translating each value by multiple of 2π .

Numerical experiments - Training




Figure: Training and validation losses in system identification training for $N_f \in \{1, 5, 10, 20\}$

Numerical experiments - Average errors


	Kin4		NeuralDyn6 with $N_f=1$		NeuralDyn6 with $N_f=5$		NeuralDyn6 with $N_f=10$		NeuralDyn6 with $N_f=20$	
	mean	std	mean	std	mean	std	mean	std	mean	std
XY	7.3967	7.8335	0.6877	1.7896	0.6205	1.6938	0.5944	1.5599	0.5718	0.7547
φ	1.1875	1.2385	0.0713	0.1840	0.0571	0.1702	0.0544	0.1532	0.0496	0.1405
v_x	0.4330	0.6145	0.1844	0.3425	0.1696	0.3263	0.1632	0.3048	0.1835	0.2965
v_y	N/A	N/A	0.0718	0.0944	0.0714	0.0932	0.0756	0.0941	0.0845	0.1059
r	N/A	N/A	0.2873	0.7396	0.2646	0.7384	0.2619	0.7375	0.2666	0.7547

Table: L^2 errors distributions of Kin4 and NeuralDyn6 by variable


Variable	XY	φ	v_x
Improvement	91.96%	95.42%	62.31%

Table: Improvements in mean error of NeuralDyn6 with $N_f = 10$ over Kin4


Numerical experiments - Average errors



(a) XY



(b) φ



(c) v_x

Figure: Average L^2 errors of Kin4 and NeuralDyn6 by stage

Numerical experiments - Open loop trajectories




Figure: Open loop trajectories of Kin4 and NODE models

Original NMPC controller

Lifted state and dynamics

$$x_k = (X_k, Y_k, \varphi_k, v_{x,k}, v_{y,k}, r_k, \delta_{k-1}), u_k = (T_k, d\delta_k)$$
$$\Rightarrow x_{k+1} = f_1(x_k, u_k) = \begin{pmatrix} \tilde{f}_1(X_k, Y_k, \varphi_k, v_{x,k}, v_{y,k}, r_k, T_k, \delta_{k-1} + d\delta_k) \\ \delta_{k-1} + d\delta_k \end{pmatrix}$$

Original NMPC controller

Lifted state and dynamics

$$x_{k+1} = f_1(x_k, u_k)$$

NMPC

$$\min_{x,u} \quad \sum_{k=0}^{N_f-1} \underbrace{\|x_k - x_k^{\text{ref}}\|_Q^2 + \|u_k\|_R^2}_{=:l(x_k, u_k)} + \underbrace{\|x_{N_f} - x_{N_f}^{\text{ref}}\|_{Q_{N_f}}^2}_{=:F(x_{N_f})}$$

$$\text{s.t. } x_0 = \hat{x}_0$$

$$x_{k+1} = f_1(x_k, u_k), \quad k = 0, \dots, N_f - 1$$

$$\underline{x} \leq x_k \leq \bar{x}, \quad k = 0, \dots, N_f$$

$$\underline{u} \leq u_k \leq \bar{u}, \quad k = 0, \dots, N_f - 1$$

Differentiable Predictive Control

Control dataset

$$\mathcal{D} = \left\{ (x_0^j, x_{0:N_f}^{\text{ref},j}) \right\}_{j=1,\dots,J}$$

DPC

$$\begin{aligned} \min_{\vartheta} \quad & \frac{1}{JN_f} \sum_{j=1}^J \sum_{k=1}^{N_f} l(x_k^j, u_k^j) + p_{S_x}(x_k^j, \underline{x}, \bar{x}) + p_{S_u}(u_k^j, \underline{u}, \bar{u}) \\ & + F(x_{N_f}^j) + p_{S_x}(x_{N_f}^j, \underline{x}, \bar{x}) \\ \text{s.t.} \quad & \left. \begin{aligned} u_{0:N_f-1}^j &= \pi_{\vartheta}(x_0^j, x_{0:N_f}^{\text{ref},j}) \\ x_{1:N_f}^j &= f_{N_f}(x_0^j, u_{0:N_f-1}^j) \end{aligned} \right\} j = 1, \dots, J \end{aligned}$$

Dataset creation

- Time-series collected at 20Hz in FSDS while running NMPC \implies transcription in the form $x_0^j, x_{0:N_f}^{\text{ref},j}$ with a sliding window
- Pre-processing: yaw signal $\varphi, \varphi^{\text{ref}}$ made continuous by appropriately translating each value by multiple of 2π .

- Time-series collected at 20Hz in FSDS while running NMPC \implies transcription in the form $x_0^j, x_{0:N_f}^{\text{ref},j}$ with a sliding window
- Pre-processing: yaw signal $\varphi, \varphi^{\text{ref}}$ made continuous by appropriately translating each value by multiple of 2π .

Numerical experiments - Constraint violation losses

Constraint	δ lower bound	δ upper bound	v_x lower bound	v_x upper bound
Violation	1.03e-5	2.14e-5	0	0

Table: Mean constraint violation committed by DPC on the validation set

Numerical experiments - Open loop behavior




Figure: Open-loop predictions of different pairs of controllers and system models.
Remark: in the subplots T , δ and $d\delta$, the blue and yellow curves are overlapping.

Numerical experiments - Closed loop behavior




Figure: Closed-loop trajectories of MPC followed by DPC in FSDS

Conclusion & Outlook

- Significant progress in the system identification, as demonstrated by the improved robustness and accuracy of NeuralDyn6 compared to Kin4.
- However, in the realm of control, we encountered challenges with the deployment of DPC in closed-loop scenarios and not enough robustness yet.
- Further improvement strategies:
 - ▶ Naive data augmentation
 - ▶ Combination of DPC with imitation learning methods such as DAgger.
⇒ use (predictive) safety filters for safe deployment and training on the real car

Conclusion & Outlook

- Significant progress in the system identification, as demonstrated by the improved robustness and accuracy of NeuralDyn6 compared to Kin4.
- However, in the realm of control, we encountered challenges with the deployment of DPC in closed-loop scenarios and not enough robustness yet.
- Further improvement strategies:
 - ▶ Naive data augmentation
 - ▶ Combination of DPC with imitation learning methods such as DAgger.
⇒ use (predictive) safety filters for safe deployment and training on the real car

Conclusion & Outlook

- Significant progress in the system identification, as demonstrated by the improved robustness and accuracy of NeuralDyn6 compared to Kin4.
- However, in the realm of control, we encountered challenges with the deployment of DPC in closed-loop scenarios and not enough robustness yet.
- Further improvement strategies:
 - ▶ Naive data augmentation
 - ▶ Combination of DPC with imitation learning methods such as DAgger.
⇒ use (predictive) safety filters for safe deployment and training on the real car

Conclusion & Outlook

- Significant progress in the system identification, as demonstrated by the improved robustness and accuracy of NeuralDyn6 compared to Kin4.
- However, in the realm of control, we encountered challenges with the deployment of DPC in closed-loop scenarios and not enough robustness yet.
- Further improvement strategies:
 - ▶ Naive data augmentation
 - ▶ Combination of DPC with imitation learning methods such as DAgger.
⇒ use (predictive) safety filters for safe deployment and training on the real car